



# Operating Systems

## Memory Management I

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>



# Programmer's Dream



Memory

- Private
- Infinitely large
- Infinitely fast
- Non-volatile
- Inexpensive


# Programmer's Dream




- Private
- Infinitely large
- Infinitely fast
- Non-volatile
- Inexpensive

Programs are getting bigger faster than memories.


# Storage/Memory Hierarchy




# Memory Hierarchy




# Memory Hierarchy



# Memory Hierarchy



# Question: Who Cares About the Memory Hierarchy?






# Memory Abstraction

- The hardware and OS memory manager make you see the memory as a single contiguous entity
- How do they do that?
  - Abstraction

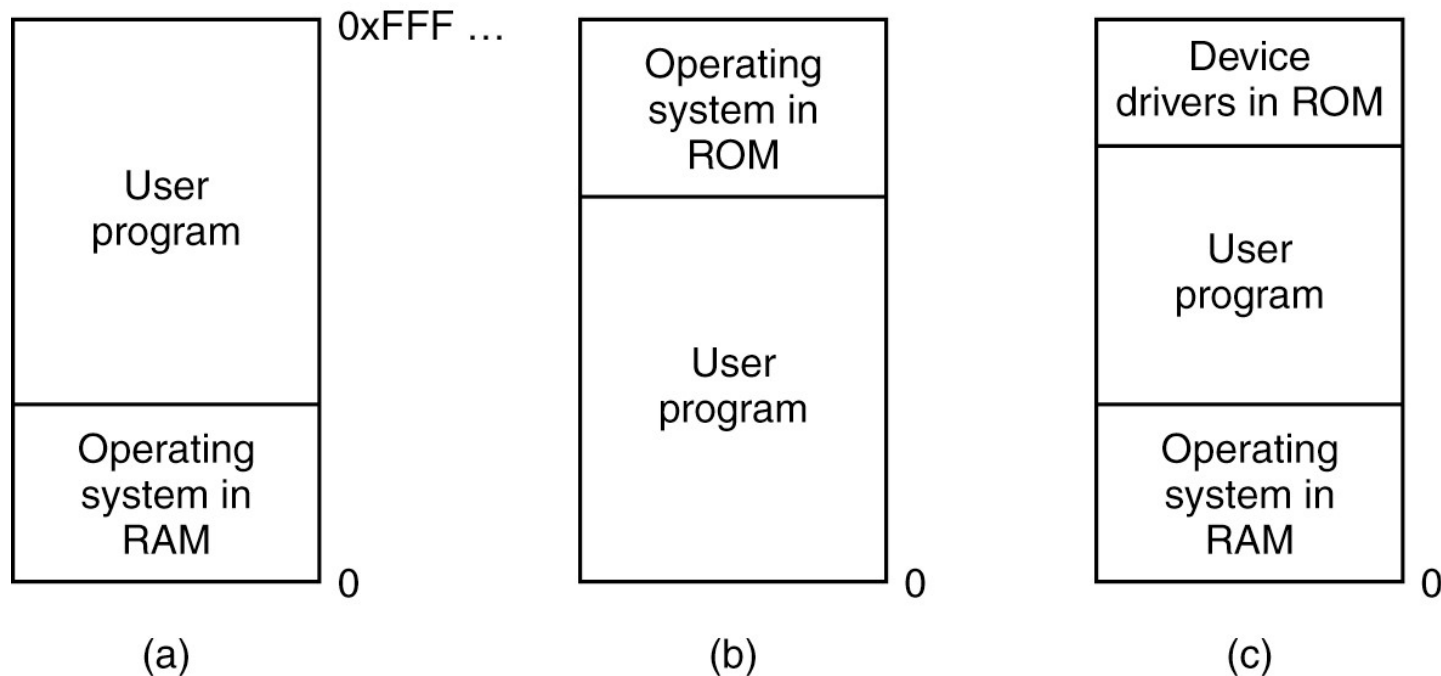
Is abstraction necessary?

# No Memory Abstraction



**Even with no abstraction, we can have several setups!**

# No Memory Abstraction



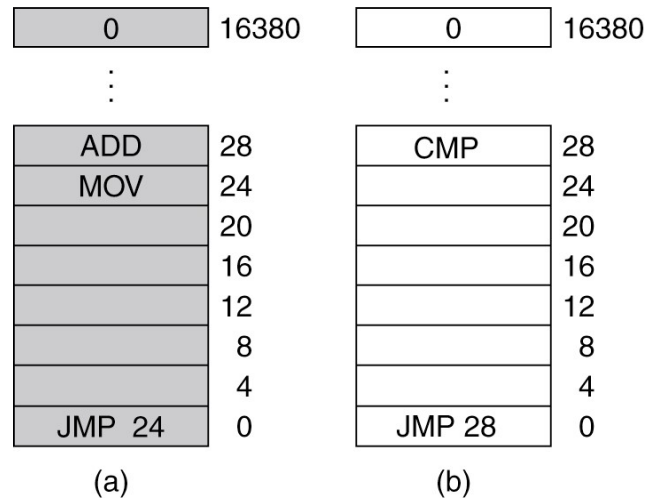
Only one process at a time can be running (threads??)

# No Memory Abstraction

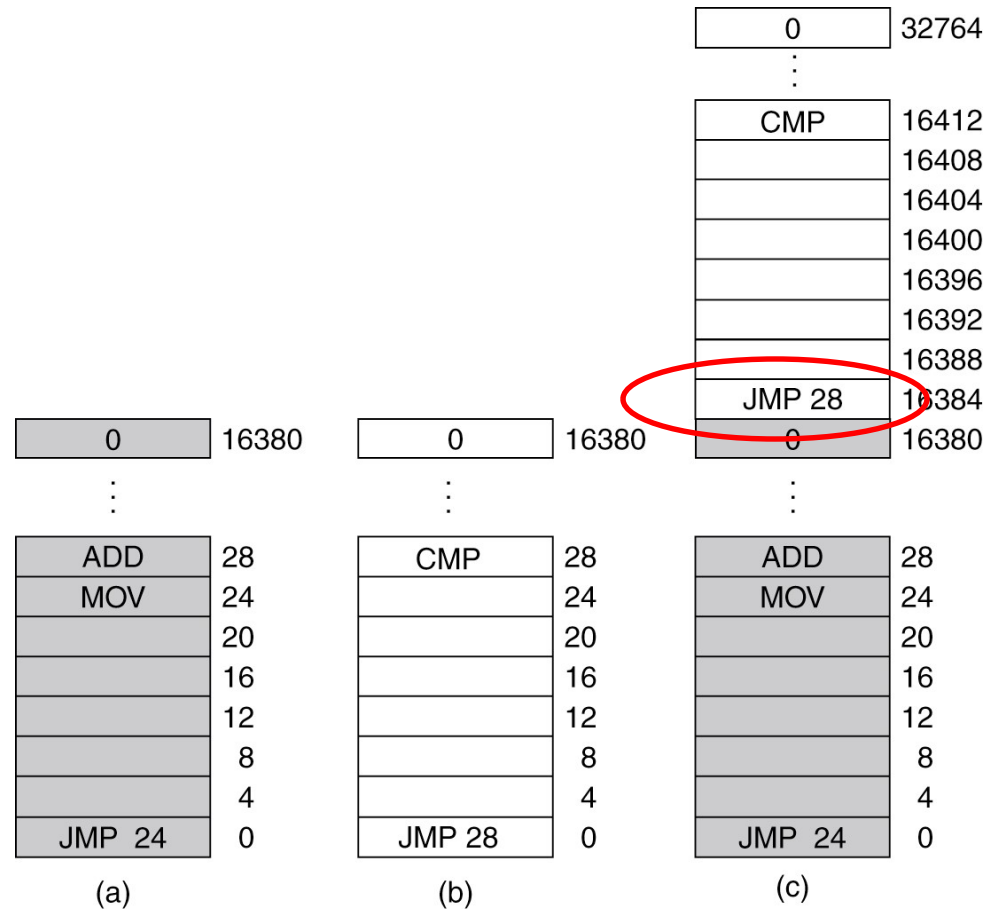
- What if we want to run multiple programs?
  - OS saves entire memory on disk
  - OS brings next program
  - OS runs next program
- We can use swapping to run multiple programs concurrently
  - Memory divided into blocks
  - Each block assigned protection bits
  - Program status word contains the same bits
  - Hardware needs to support this
  - Example: IBM 360

Swapping

# No Memory Abstraction



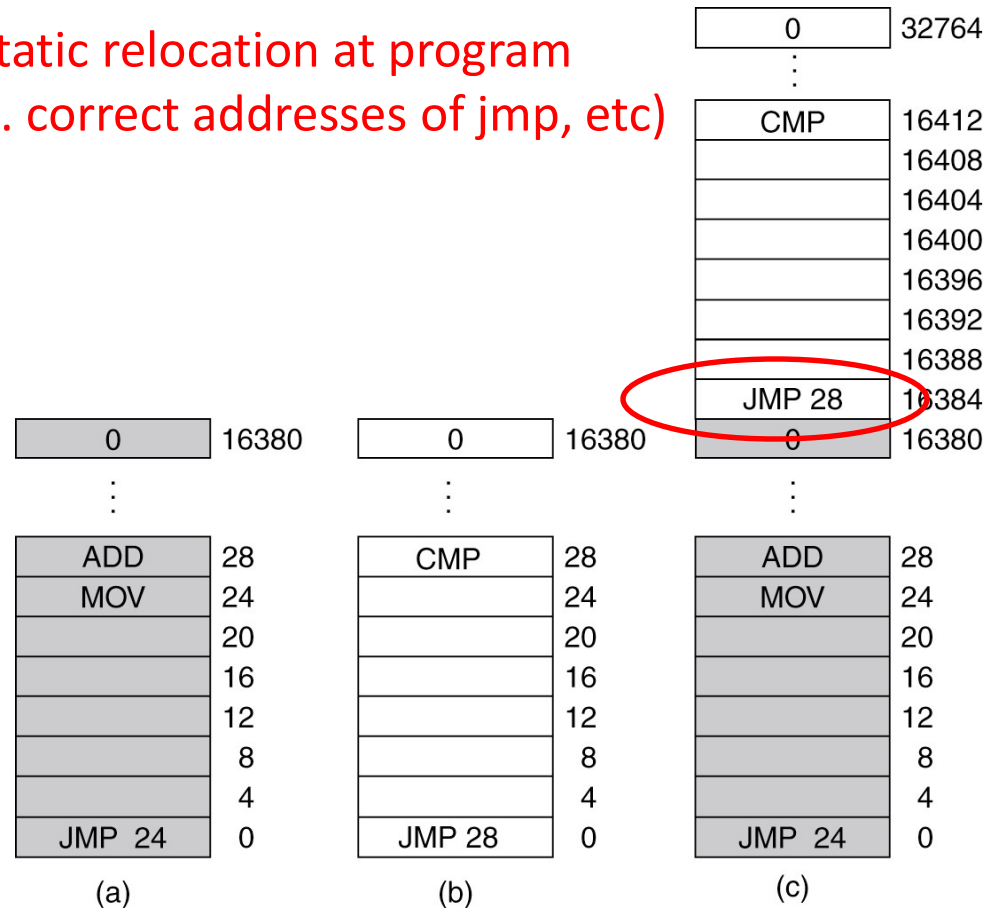
# No Memory Abstraction



Using absolute address is wrong here

# No Memory Abstraction

We can use static relocation at program load time (i.e. correct addresses of jmp, etc)



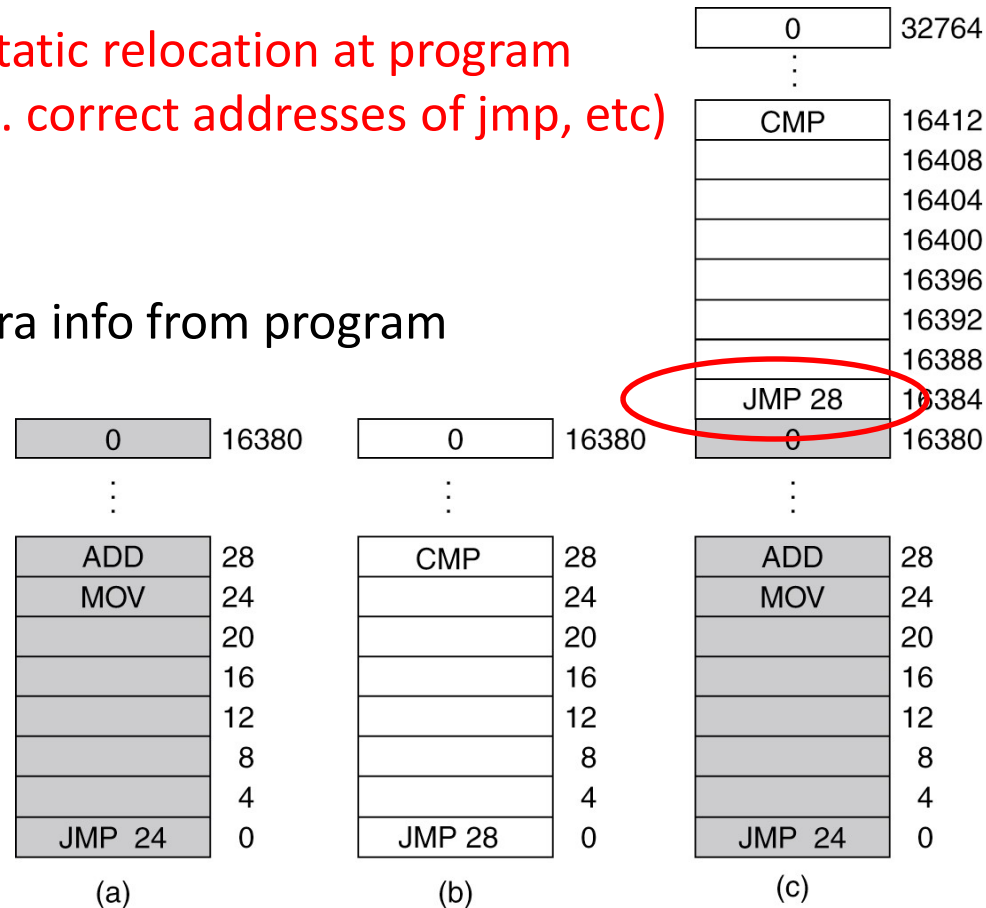
Using absolute address is wrong here

# No Memory Abstraction

We can use static relocation at program load time (i.e. correct addresses of jmp, etc)

**Bad Idea!**

- Slow
- Require extra info from program



Using absolute address is wrong here

Bottom line: Memory abstraction is needed!



# Memory Abstraction

- To allow several programs to co-exist in memory we need
  - Protection
  - Relocation
  - Sharing
  - Logical organization
  - Physical organization
- A new abstraction for memory: **Address Space**
- Address space = set of addresses that a process can use to address memory

# Protection

- Processes need to acquire permission to reference memory locations for reading or writing purposes.
- Location of a program in main memory is unpredictable.
- Memory references generated by a process must be checked at run time.

# Relocation

- Programmers typically do not know in advance which other programs will be resident in main memory at the time of execution of their program.
- Active processes need to be able to be swapped in and out of main memory in order to maximize processor utilization.
- Specifying that a process must be placed in the same memory region when it is swapped back in would be limiting
  - may need to *relocate* the process to a different area of memory

# Sharing

- It is advantageous to allow each process access to the same copy of the program/library/... rather than have their own separate copy.
- Memory management must allow controlled access to shared areas of memory without compromising protection

# Logical Organization

- We see memory as linear one-dimensional address space.
- A program = code + data + ... = modules
- Those modules must be organized in that logical address space

# Physical Organization

- Memory is really a hierarchy
  - Several levels of caches
  - Main memory
  - Disk
- Managing the different modules of different programs in such a way as:
  - To give illusion of the logical organization
  - To make the best use of the above hierarchy

All of this must be done while ensuring:

- **Transparency:** the running programs must not know that all of this is happening.
- **Efficiency:** both in terms of time (speed) and space (not wasting a lot of memory)
- **Protection:** as we saw, protecting processes from each other

# Address Space: Base and Limit

- Map each process address space onto a different part of physical memory
- Two registers: Base and Limit
  - **Base**: start address of a program in physical memory
  - **Limit** (sometimes called **bound**): length of the program
- For every memory access
  - Base is added to the address
  - Result compared to Limit
  - Who is doing this? A piece of hardware inside the processor called the **memory management unit (MMU)**.
- Only OS can modify Base and Limit



# Address Space: Base and Limit

**Main drawbacks:**

Need to add and compare for each  
memory address


What if memory space is not enough for  
all programs?

# Address Space: Base and Limit

## Main drawbacks:


Need to add and compare for each  
memory address

What if memory space is not enough for  
all programs?




The we may need to **swap** some programs out of the memory.

# Swapping

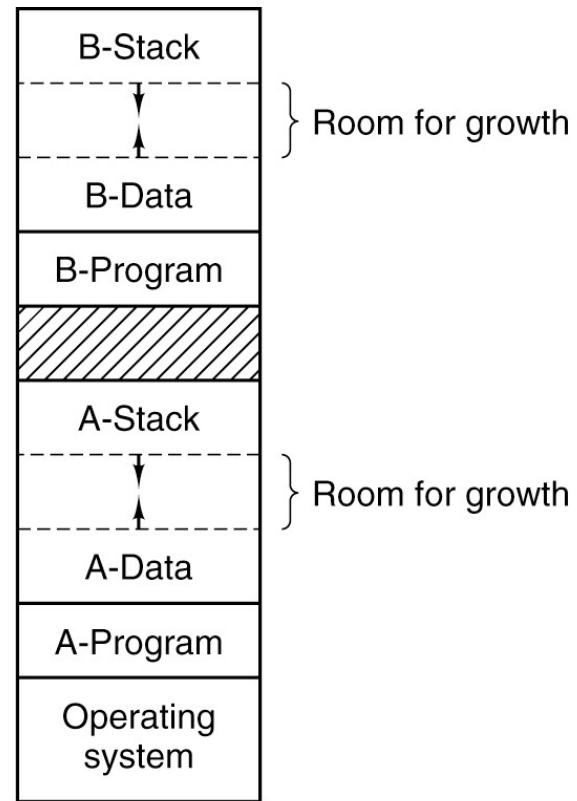


# Swapping

- Programs move in and out of memory
- **Holes** are created
- Holes can be combined -> **memory compaction**
- What if a process needs more memory?
  - If a hole is adjacent to the process, it is allocated to it
  - Process must be moved to a bigger hole
  - Process suspended till enough memory is there



(a)

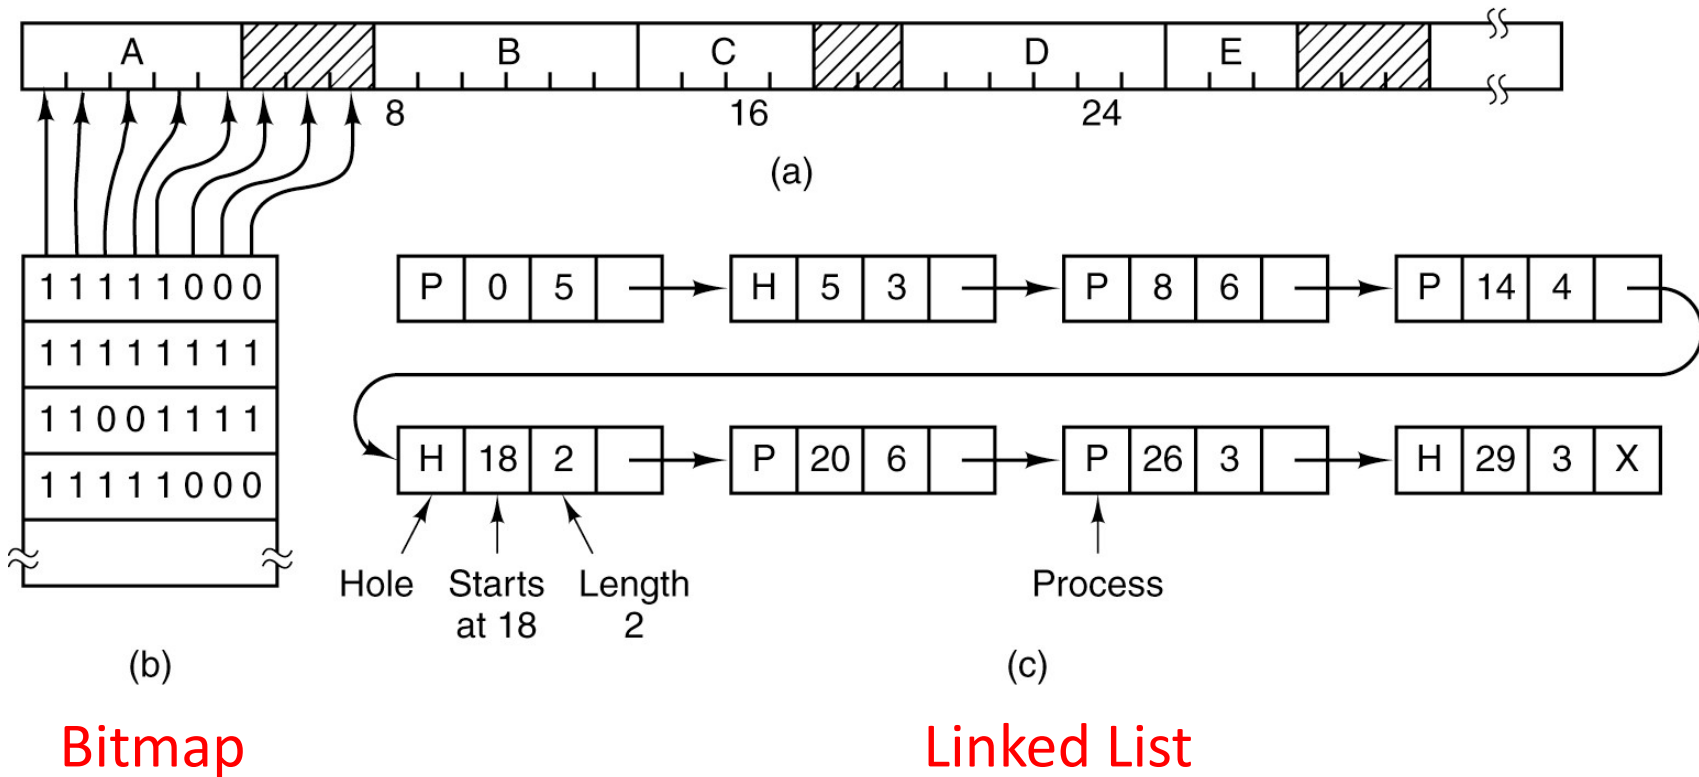


(b)

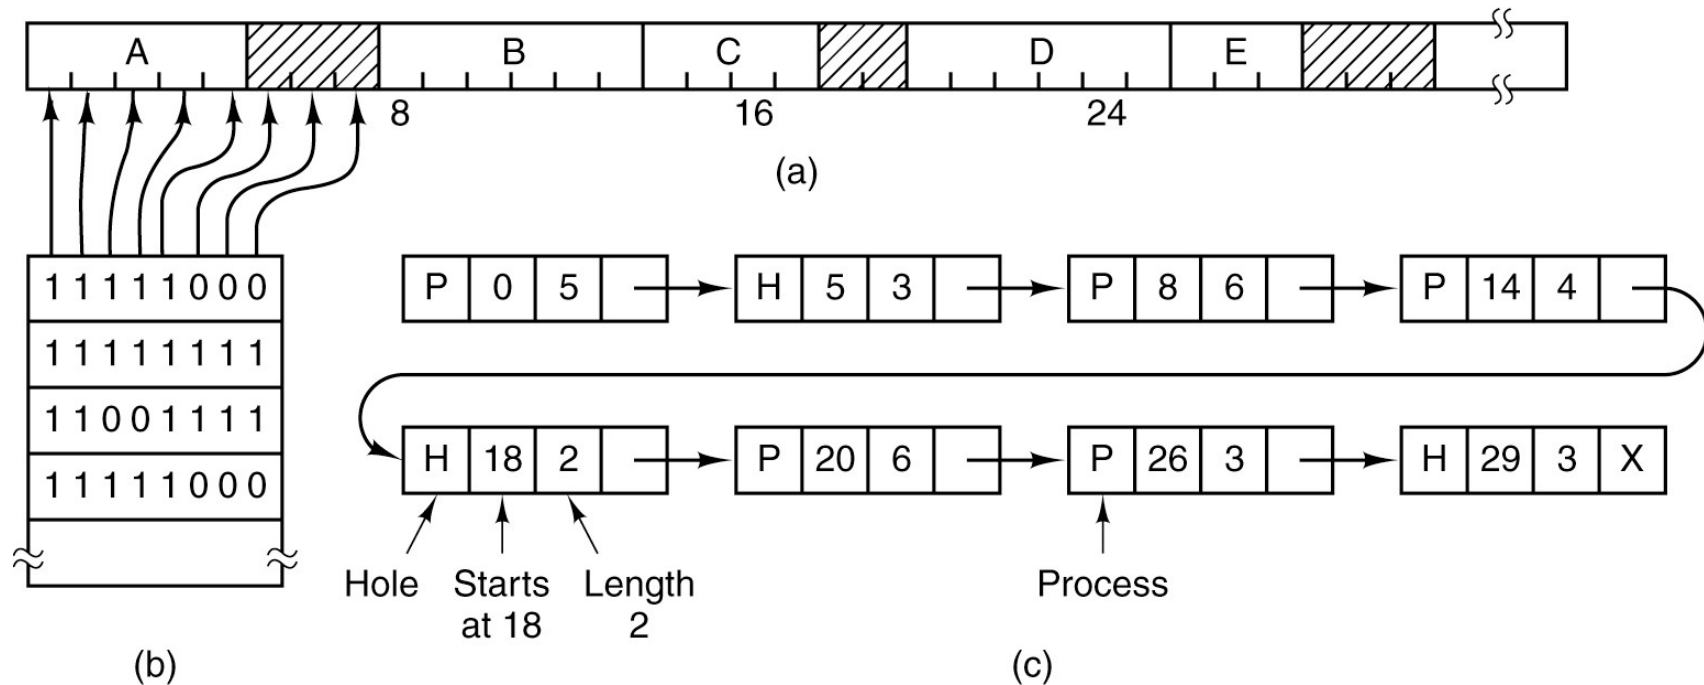
# Managing Free Memory

- Bitmap:
  - Memory is divided into allocation **units of equal size**.
  - Each unit has a corresponding bit in the bitmap.
  - 0 = unit is free    1 = unit is occupied (or vice versa, doesn't matter).
- Linked List: of allocated and free memory segments.
  - **Segments are of different sizes.**

# Managing Free Memory



# Managing Free Memory



Bitmap




Slow: To find k-consecutive 0s for a new process

Linked List



# Managing Free Memory: Linked List

- Linked list of allocated and free memory **segments**
- More convenient be double-linked list



# Managing Free Memory: Linked List

- How to allocate?
  - First fit
  - Best fit
  - Next fit
  - Worst fit
  - ...

# Memory Management Techniques

- Memory management brings processes into main memory for execution by the processor
  - involves **virtual memory**
  - based on:
    - **segmentation** (variable size parts) or
    - **paging** (fixed size parts) ← next lecture

# Conclusions

- Process is CPU abstraction
- Address space is memory abstraction
  - OS memory manager and the hardware helps providing this abstraction
- Two main tasks needed from OS regarding memory management:
  - managing free space
  - making best use of the memory hierarchy